

VMware Tanzu RabbitMQ

An Overview

Dan Buchko, VMware Tanzu Solutions Engineer
July 16, 2020

Introduction

What is RabbitMQ?

- A long-running open-source project written in Erlang backed by VMware
- A general-purpose message broker designed with a smart broker / passive consumer model
- Enables cross-language messaging by implementing Advanced Message Queueing Protocol specification
- Support for range of additional protocols (STOMP, MQTT, AMQP, etc.)
- A tool with wide and extendable support (many client libraries, many plugins)

Why RabbitMQ? [1/2]

- Most widely used & deployed open-source messaging tool with a proven track record
- Popular with a strong community that creates resources to extend and support its capabilities (e.g., community plug-ins)
- Built off of Erlang for high availability
- Pluggable authentication support (LDAP)
- Empowers developers with developer APIs for all actions (incl config)

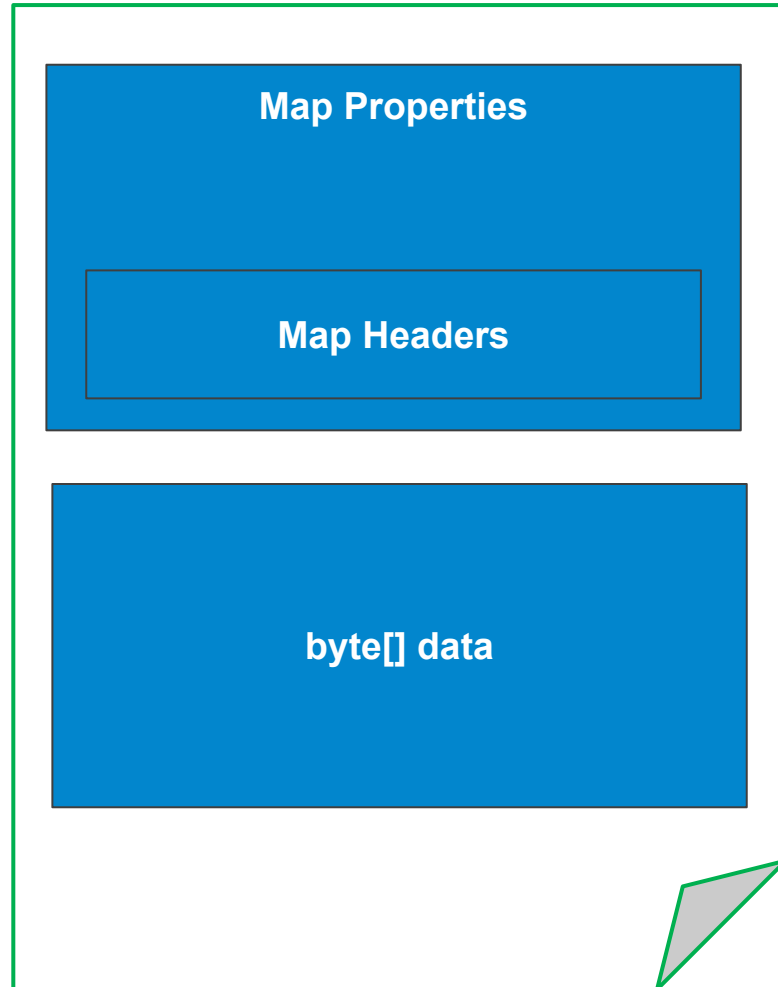
Why RabbitMQ? [2/2]

- Routing from smart broker unlocks powerful messaging capabilities
- Flexibility for almost any language to share messages across operating systems and environments
- Get up & running in <20 minutes
- Extendibility in plug-ins and protocols means it can be customized for wide range of needs

RabbitMQ Concepts

RabbitMQ / AMQP Terminology

Messages



Properties/headers are just Key-Value pairs

Immutable properties: message-id, user-id ...
routing-key
delivery-mode (2 = persistent)
content-type
reply-to
correlation-id

Application-specific properties


Data is opaque (just an array of bytes)

-> AMQP is language agnostic

RabbitMQ / AMQP Terminology

Producer / Publisher


Applications that create messages and deliver them to RabbitMQ Broker.



RabbitMQ / AMQP Terminology

Consumer / Subscriber

Applications that receive messages from the broker



RabbitMQ / AMQP Terminology

Exchanges

Exchanges reside on the broker

Can be:

- Durable (persisted to disk)
- Transient (in memory)

Publishers always publish to an exchange


Responsible for routing messages based on type and message content

Stateless

Bound to queues

4 main types:

- Direct
- Fanout
- Topic
- Headers




RabbitMQ / AMQP Terminology

Bindings

Define associations between resources, eg

- Exchanges bound to queues
- Exchanges bound to other exchanges



RabbitMQ / AMQP Terminology

Queues

Store messages, i.e. stateful

Consumers always consume from queues

Can be configured to have range of behavior:

- Durable (persisted to disk) / transient

- Expiration


- Lazy

- HA

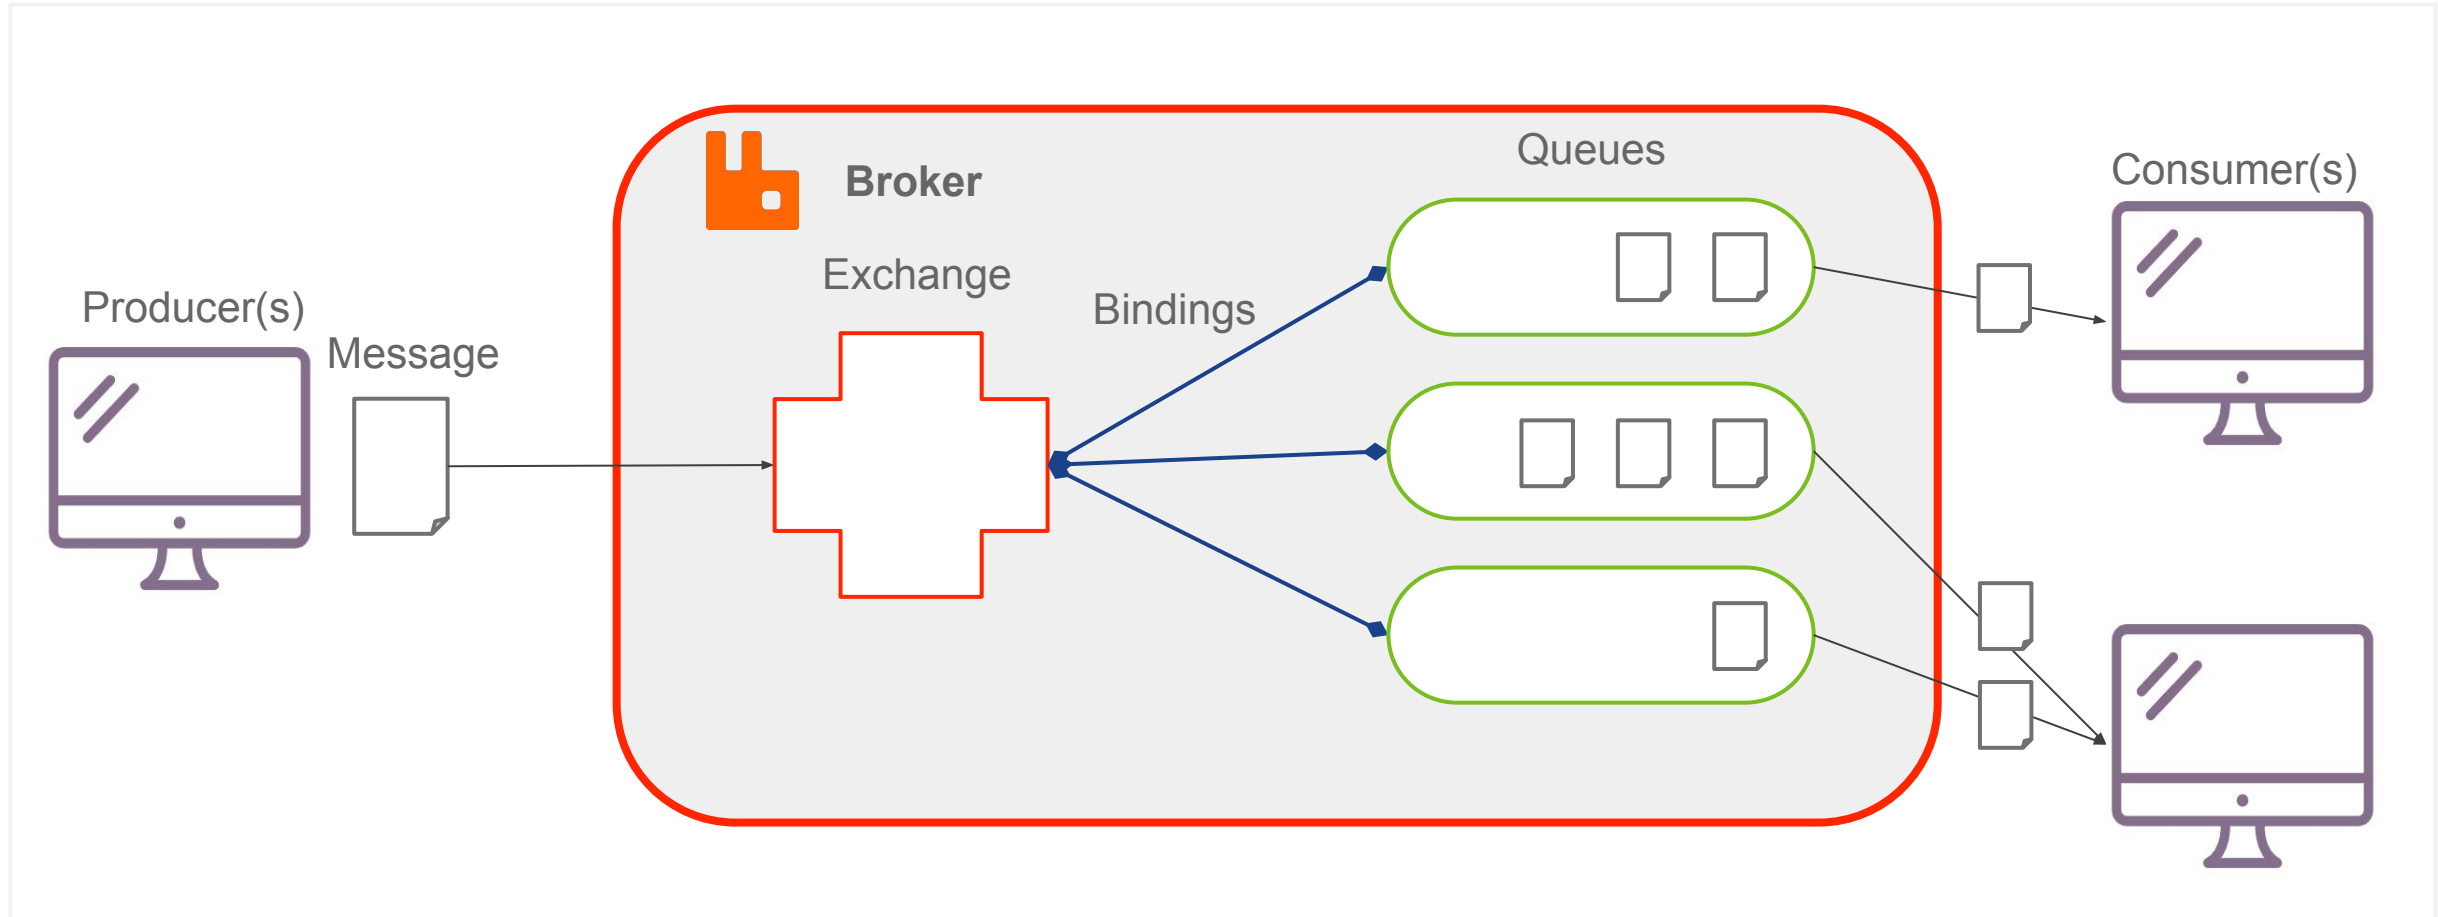
- Quorum

- etc.

HA and Quorum queues give control over reliability and performance.



Putting it All Together



Clustering

Clustering

Introduction

A Cluster is two or more nodes working together to form one logical broker

Clustering brings a number of benefits:


Improved performance (* if used wisely)

Multiple queues across multiple machines

Higher availability

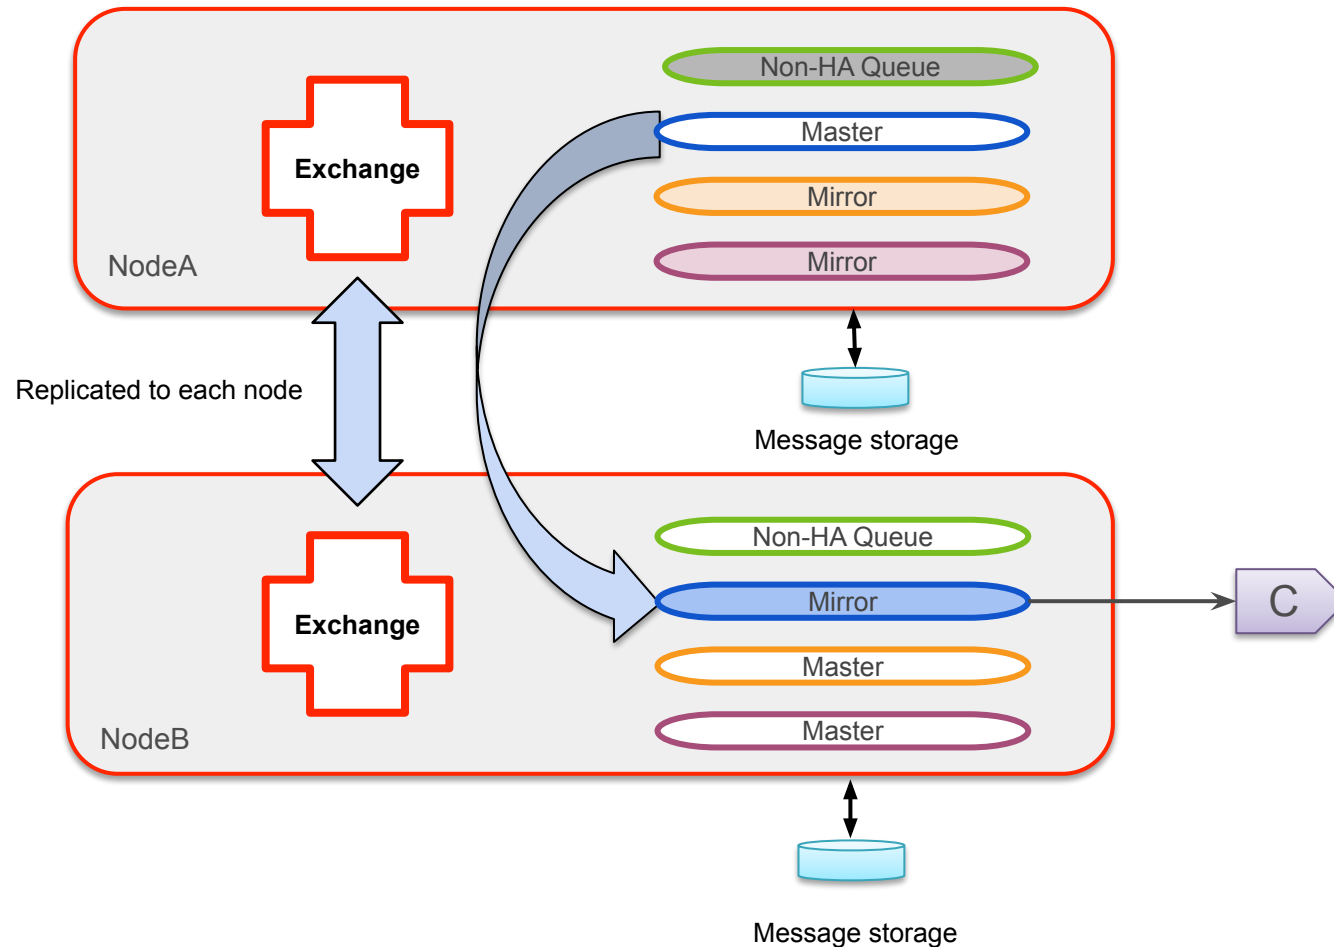
Higher reliability

Mirrored queues



A 2-Node Cluster

HA vs Non-HA Queues



Non-HA Queue

- All messages are held on the node where it was created
- Other nodes act as proxies for clients which connect to them

HA-Queue

- Client connects to a node and creates a queue. The "master" for this queue is now on this node
- Mirror queues are created on other nodes in the cluster - how many is configurable by policy
- Mirrors hold copies of all messages
- Distribution of messages to consumers is managed by the master node - so there is always network traffic back to the master
- If the node holding the master queue fails, remaining mirrors hold an election and the winner becomes the new master

High Availability

What Happens If A Node Fails?

What happens if one node in a cluster fails?

- All messages in queues which reside on that node are lost unless messages are persistent
- Persistent messages can only be recovered if the failed node can be recovered

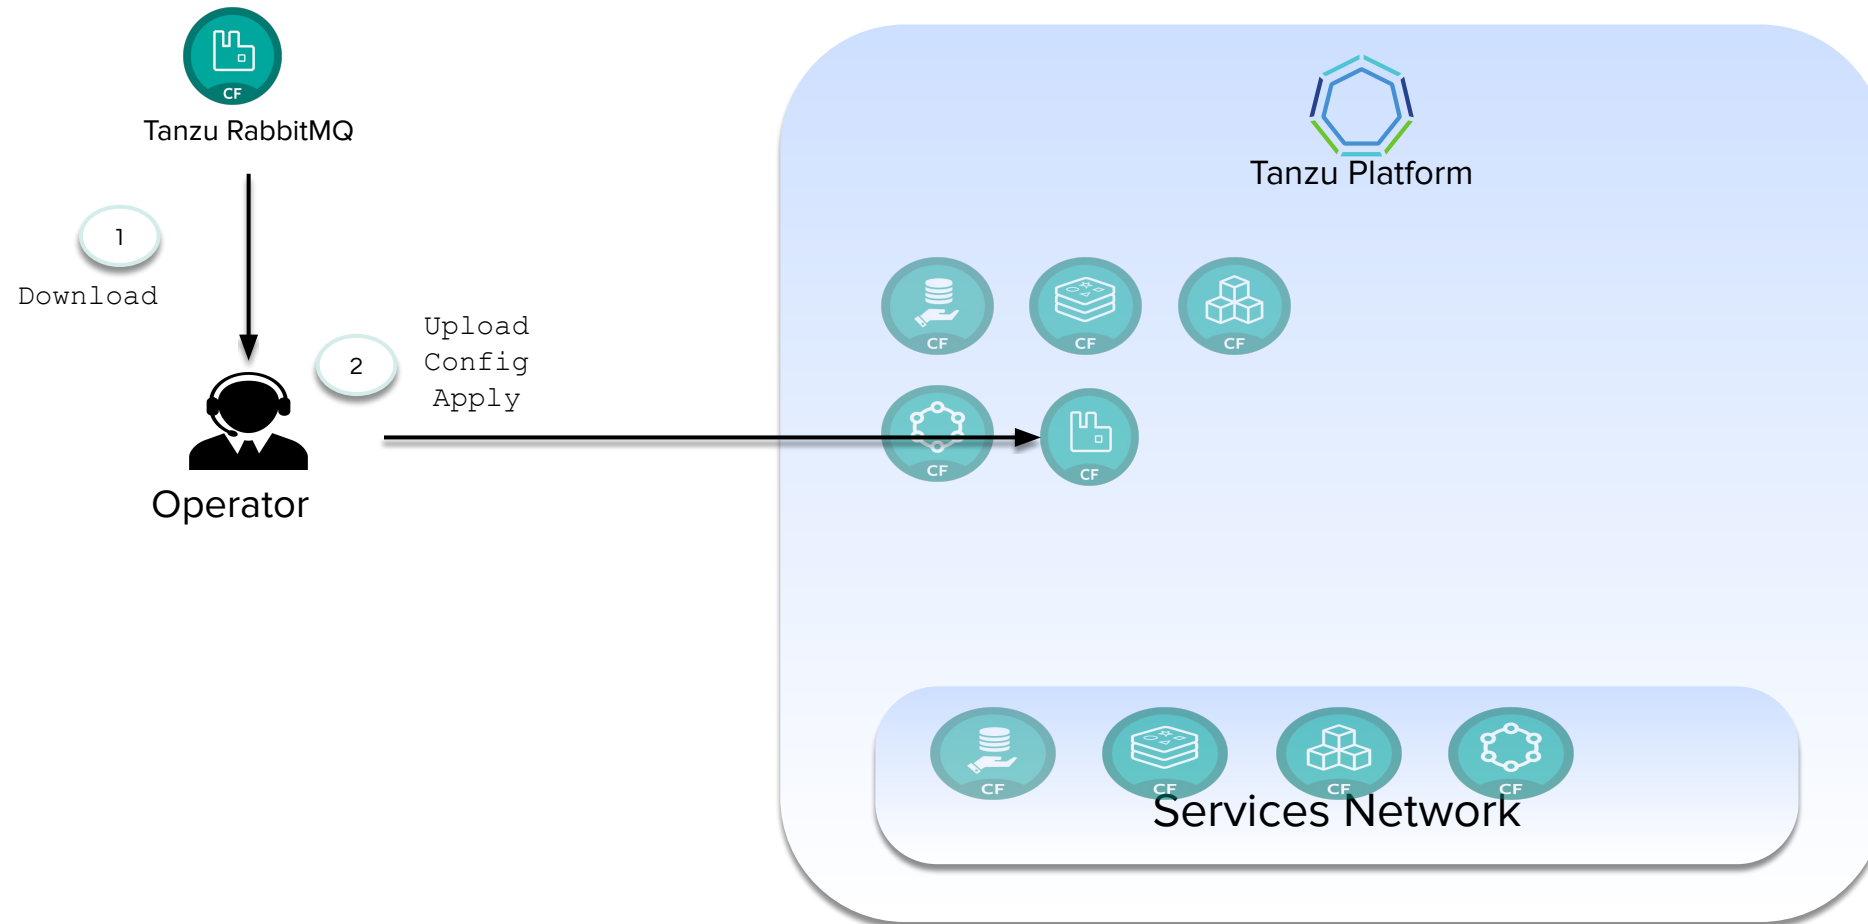
Solution – HA

- Replicates messages across other nodes in the cluster
- If a node dies, other nodes have copies of all messages
- Clients reconnect to the cluster and continue where they left off
- Configured by policy on a per-queue basis (`ha-mode={all, exactly, nodes}`)
- Transactions and Publisher Confirms span all mirrors

Tanzu RabbitMQ (RabbitMQ on PCF)

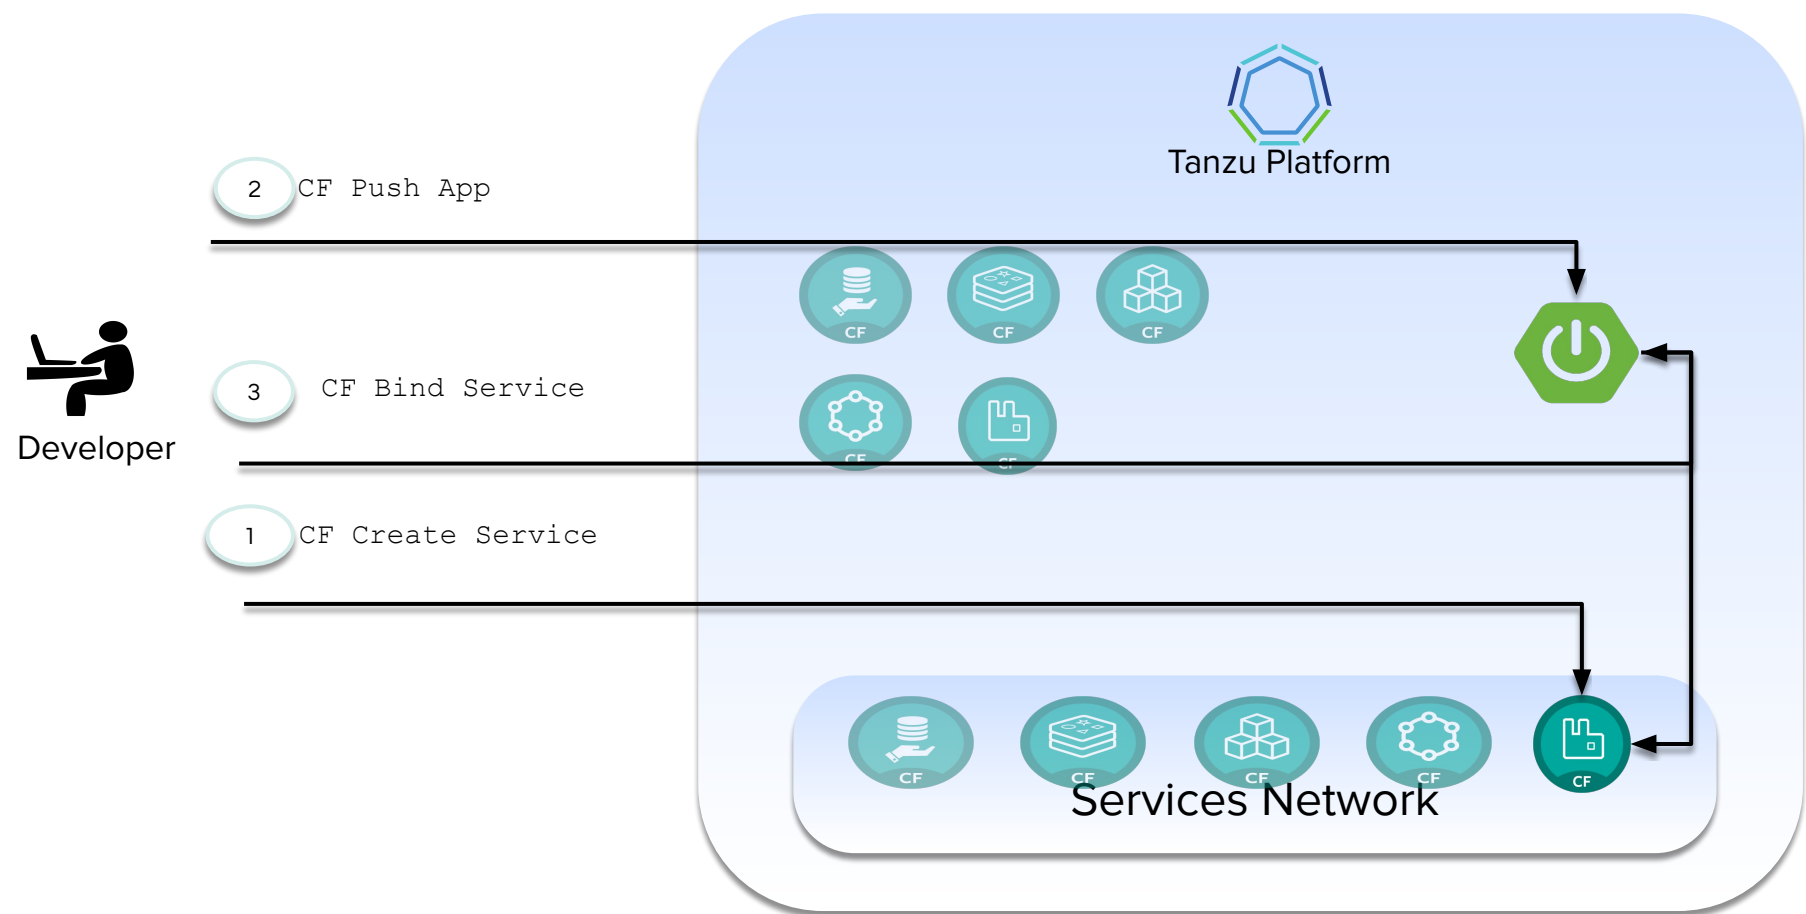
RabbitMQ on Tanzu Platform

Operator Flow



RabbitMQ on Tanzu Platform

Developer Flow



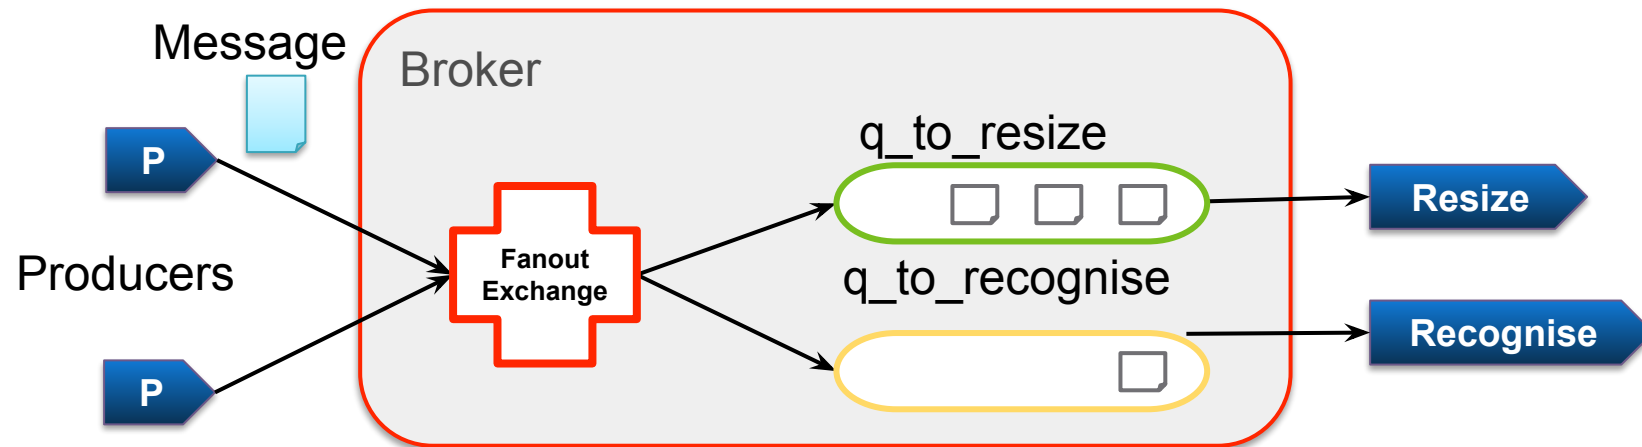
Thank You!

Appendix: Exchange Types



Fanout Exchange

Publish/Subscribe



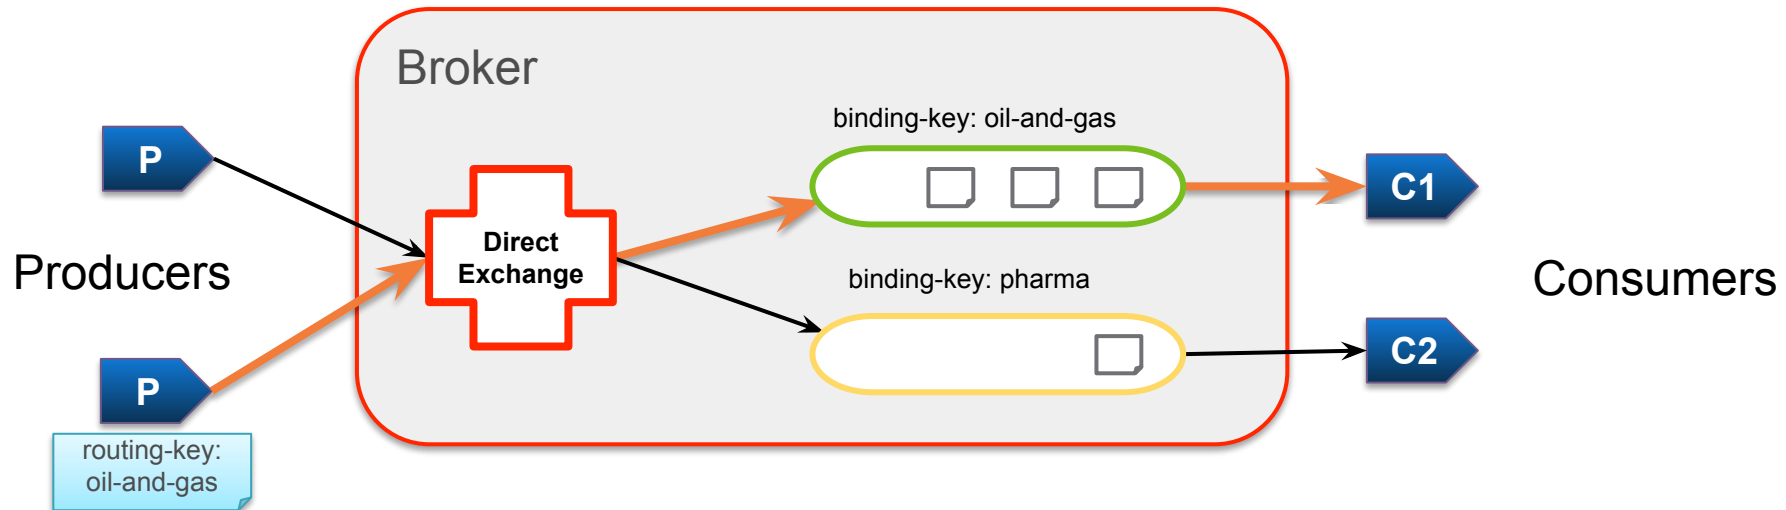
Every queue gets a copy of every message

Upload a photo to
Facebook.

Consumers resizing,
doing face recognition...

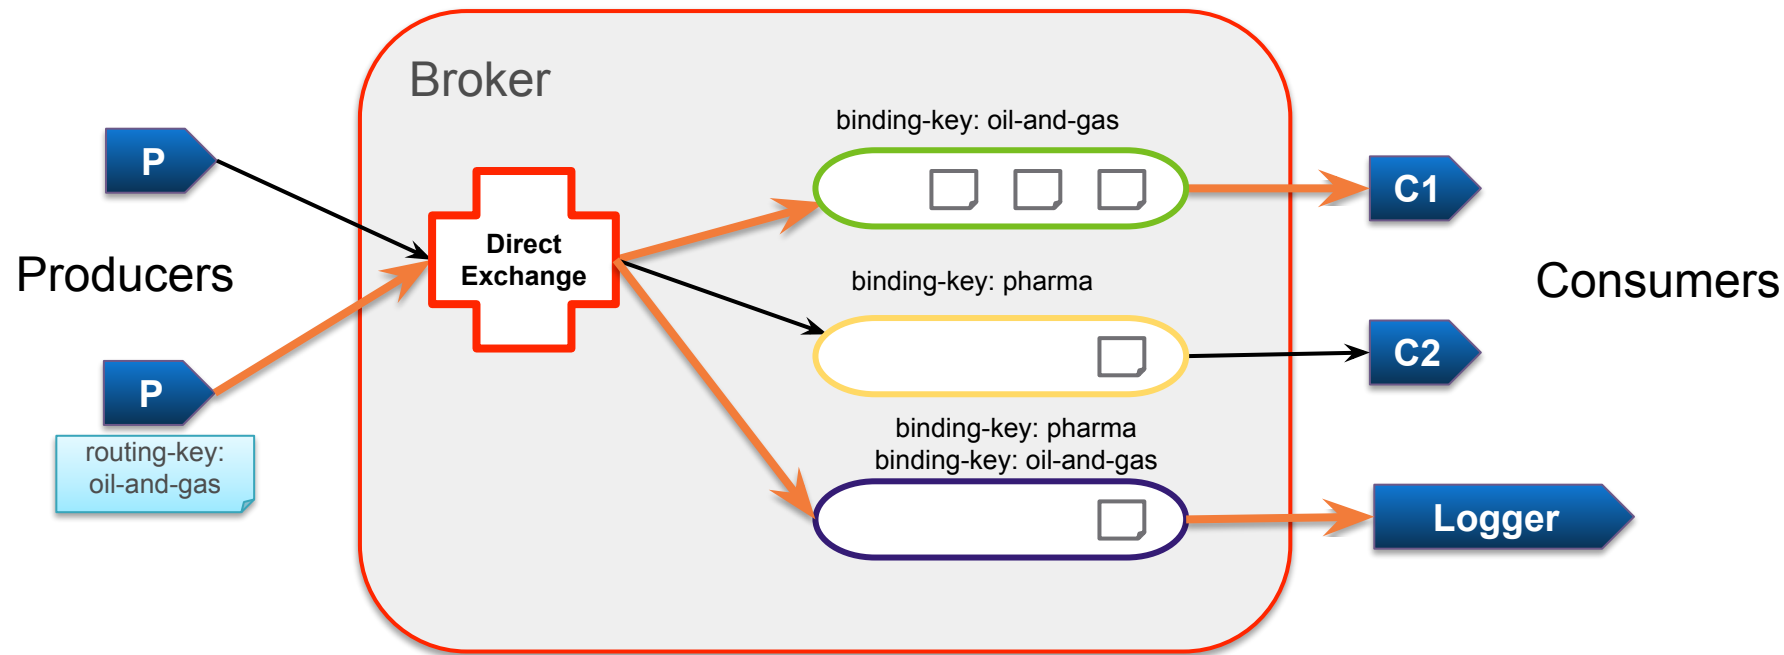
Direct Exchange

Routing – The Direct Exchange



- Queues 'bind' to the exchange
- A queue can have more than one binding-key
- A message is routed to each queue whose 'binding-key' matches the message's 'routing-key'

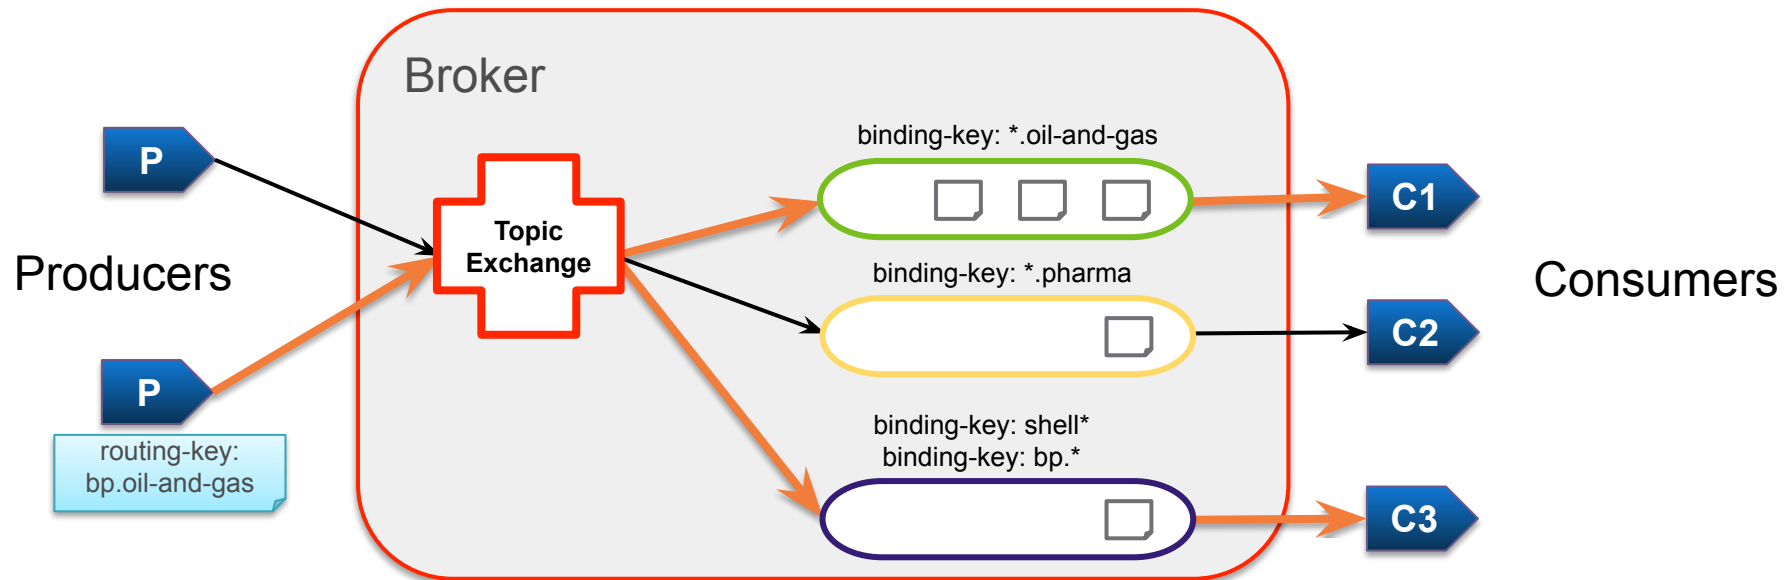
Direct Exchange - Multiple Bindings



Different queues can use the same binding key

Topic Exchange

Pattern Matching The Routing-Key

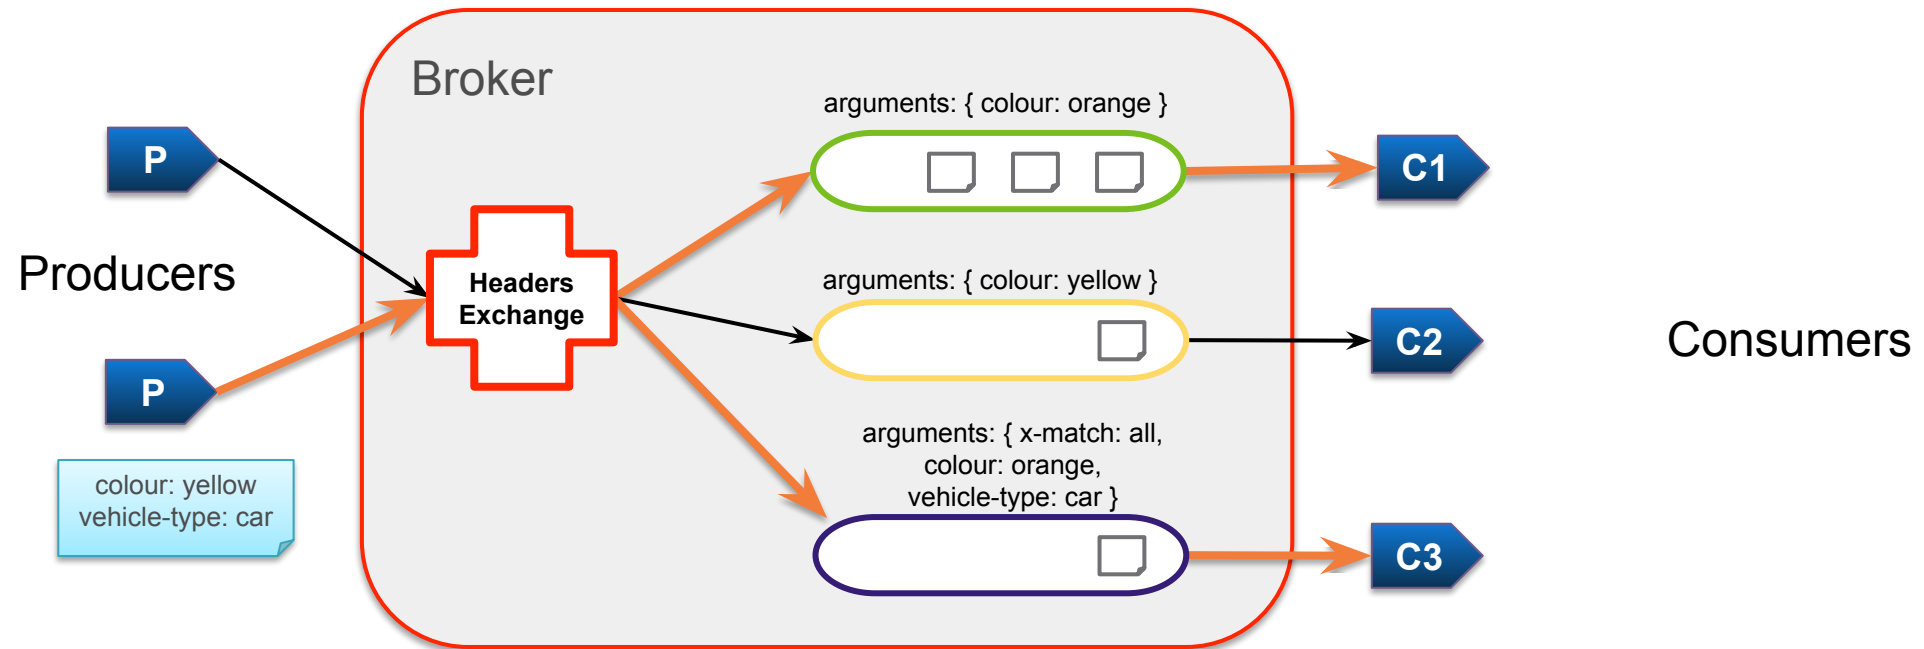


Bindings match word patterns
Words are separated by dots '.'

- zero or more words
* - exactly one word

Headers Exchange

Not Just The Routing Key



Binding arguments specify a list of headers to match exactly
Optional x-match argument: 'all' or 'any' gives AND/OR